Manage a PB/day of artifacts without breaking a sweat (or your infrastructure)

Jenkins/DevOps World

Anupam Garg
Kyle Harris

Warning

This is supposed to be a “lightning” talk, but I’m from the south …

Whole lotta data

Our build/test cycles generate a lot of data

Whole lotta data

Our build/test cycles generate a lot of data

115 GB of build artifacts (product + toolchain + test binaries)

Whole lotta data

Our build/test cycles generate a lot of data

115 GB of build artifacts (product + toolchain + test binaries)
500 slaves provisioned daily to run tests (each requiring a copy)

Whole lotta data

Our build/test cycles generate a lot of data

115 GB of build artifacts (product + toolchain + test binaries)
500 slaves provisioned daily to run tests (each requiring a copy)
6000 desktop tests (single host testbed)

Whole lotta data

Our build/test cycles generate a lot of data

115 GB of build artifacts (product + toolchain + test binaries)
500 slaves provisioned daily to run tests (each requiring a copy)
6000 desktop tests (single host testbed)
1100 virtual tests (multple hosts testbed)

Traditional copy artifact workflow

(hop through master not shown)

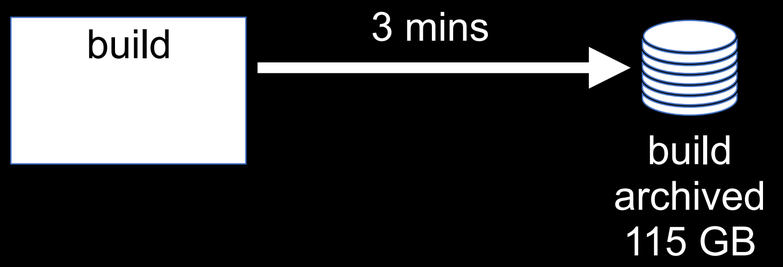

Traditional copy artifact workflow

(hop through master not shown)

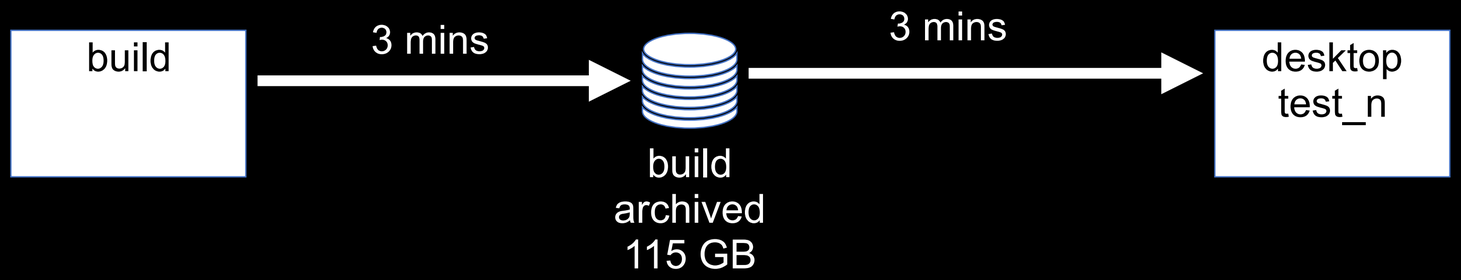

Traditional copy artifact workflow

(hop through master not shown)

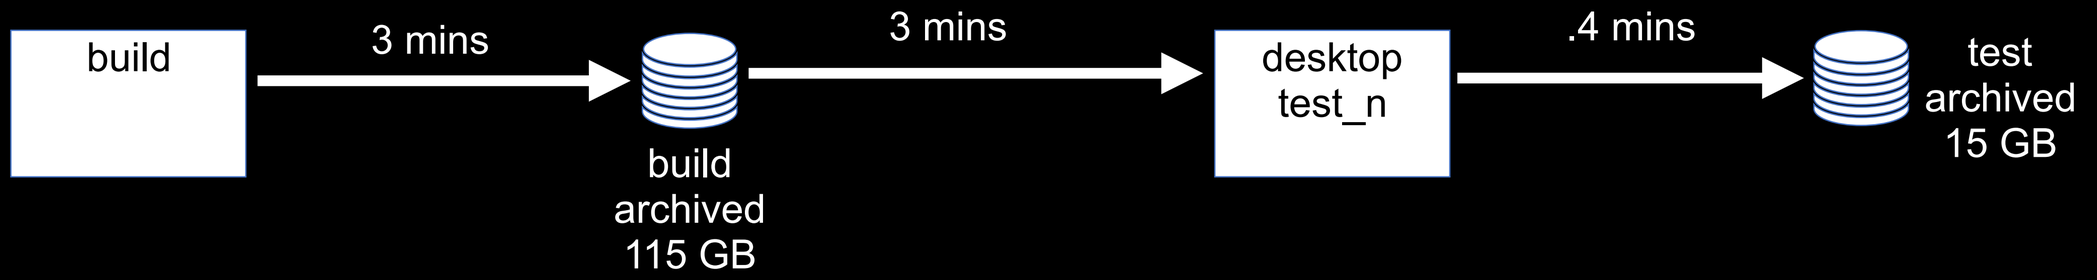

Traditional copy artifact workflow

(hop through master not shown)

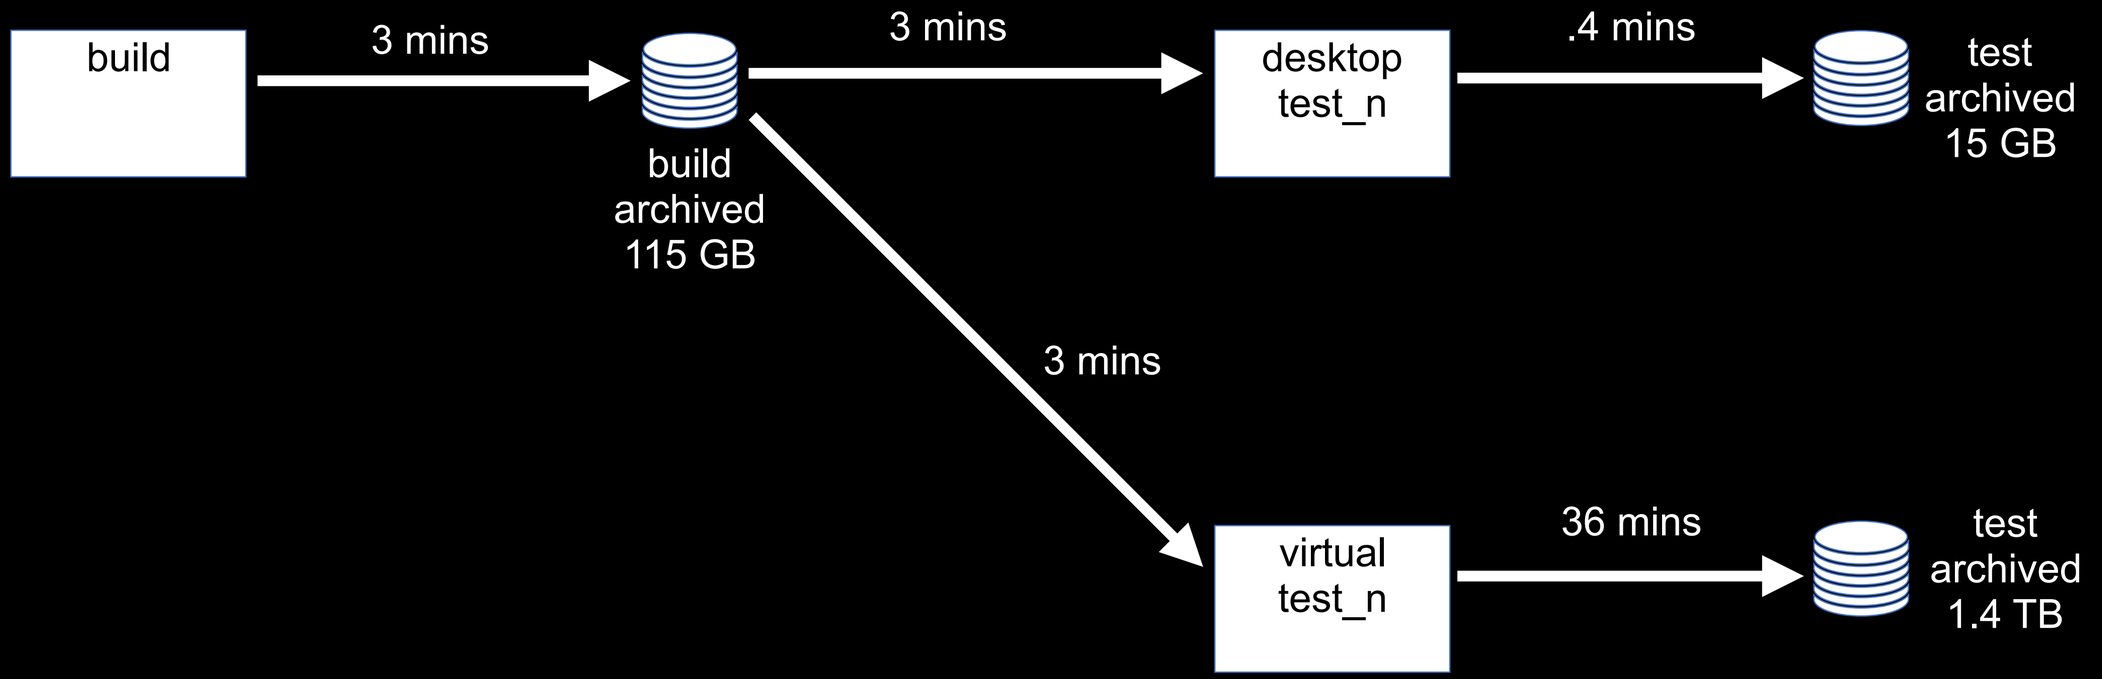

Traditional copy artifact workflow

(hop through master not shown)

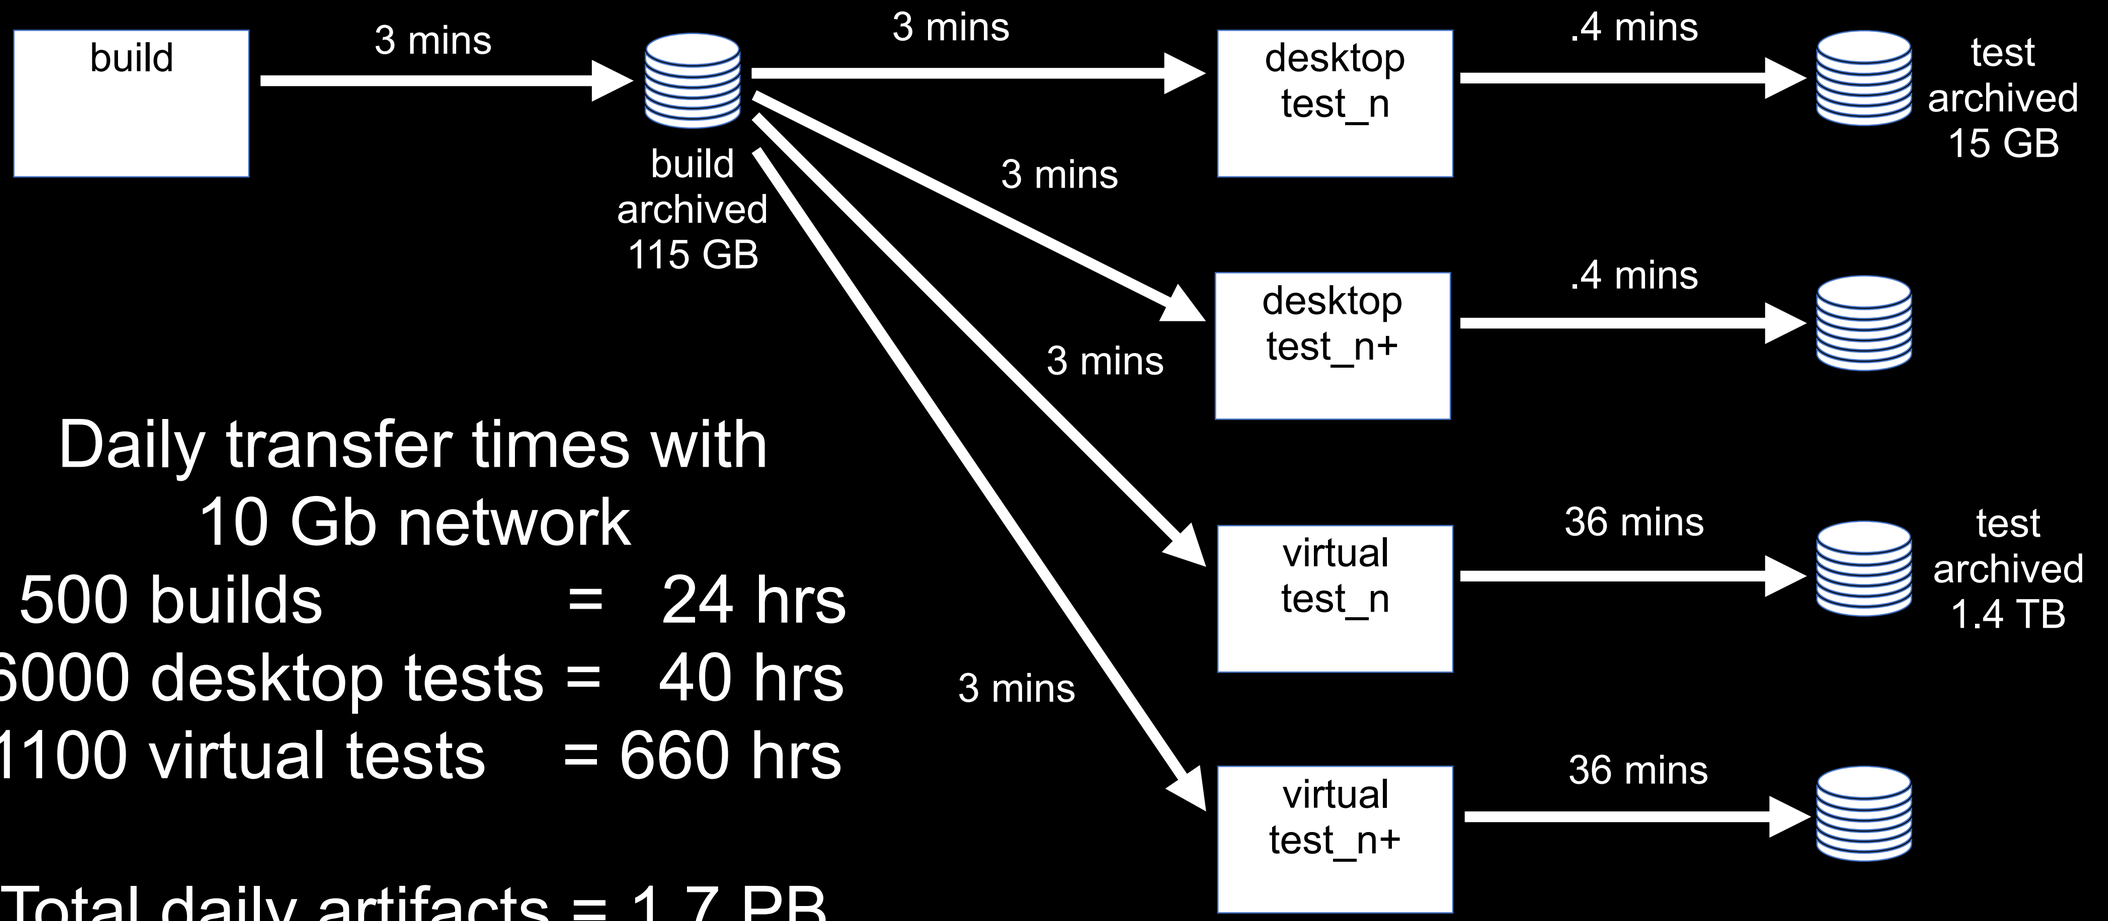

Daily transfer times with
10 Gb network

500 builds = 24 hrs
6000 desktop tests = 40 hrs
1100 virtual tests = 660 hrs

Total daily artifacts = 1.7 PB

Zero copy cloning workflow

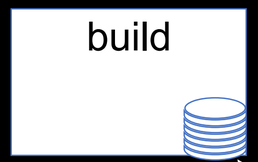

artifacts are contained within the slave
the slave is saved and cloned for downstream builds

Zero copy cloning workflow

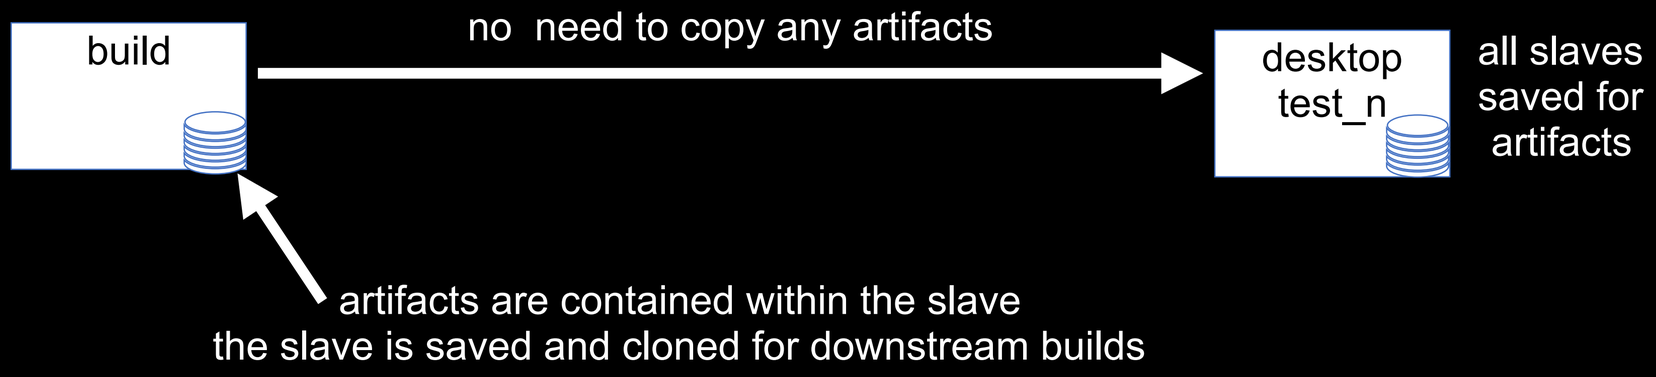

Zero copy cloning workflow

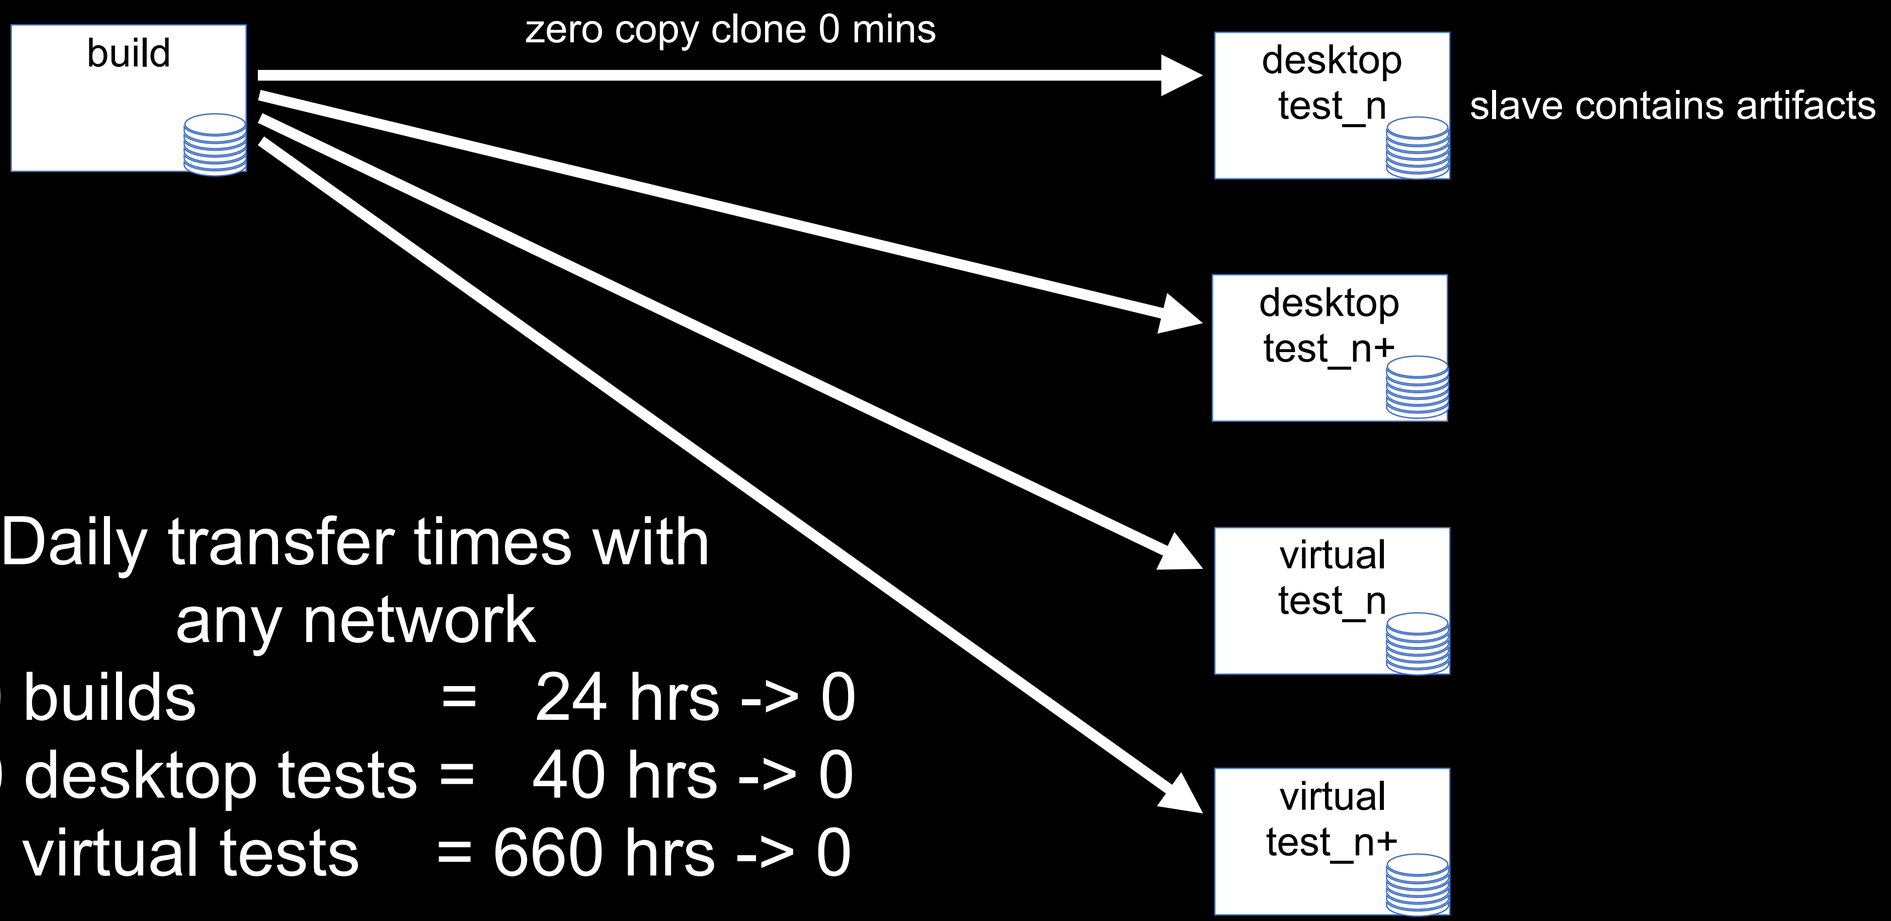

Daily transfer times with
any network

500 builds = 24 hrs -> 0
6000 desktop tests = 40 hrs -> 0
1100 virtual tests = 660 hrs -> 0

Incremental to monolithic artifacts

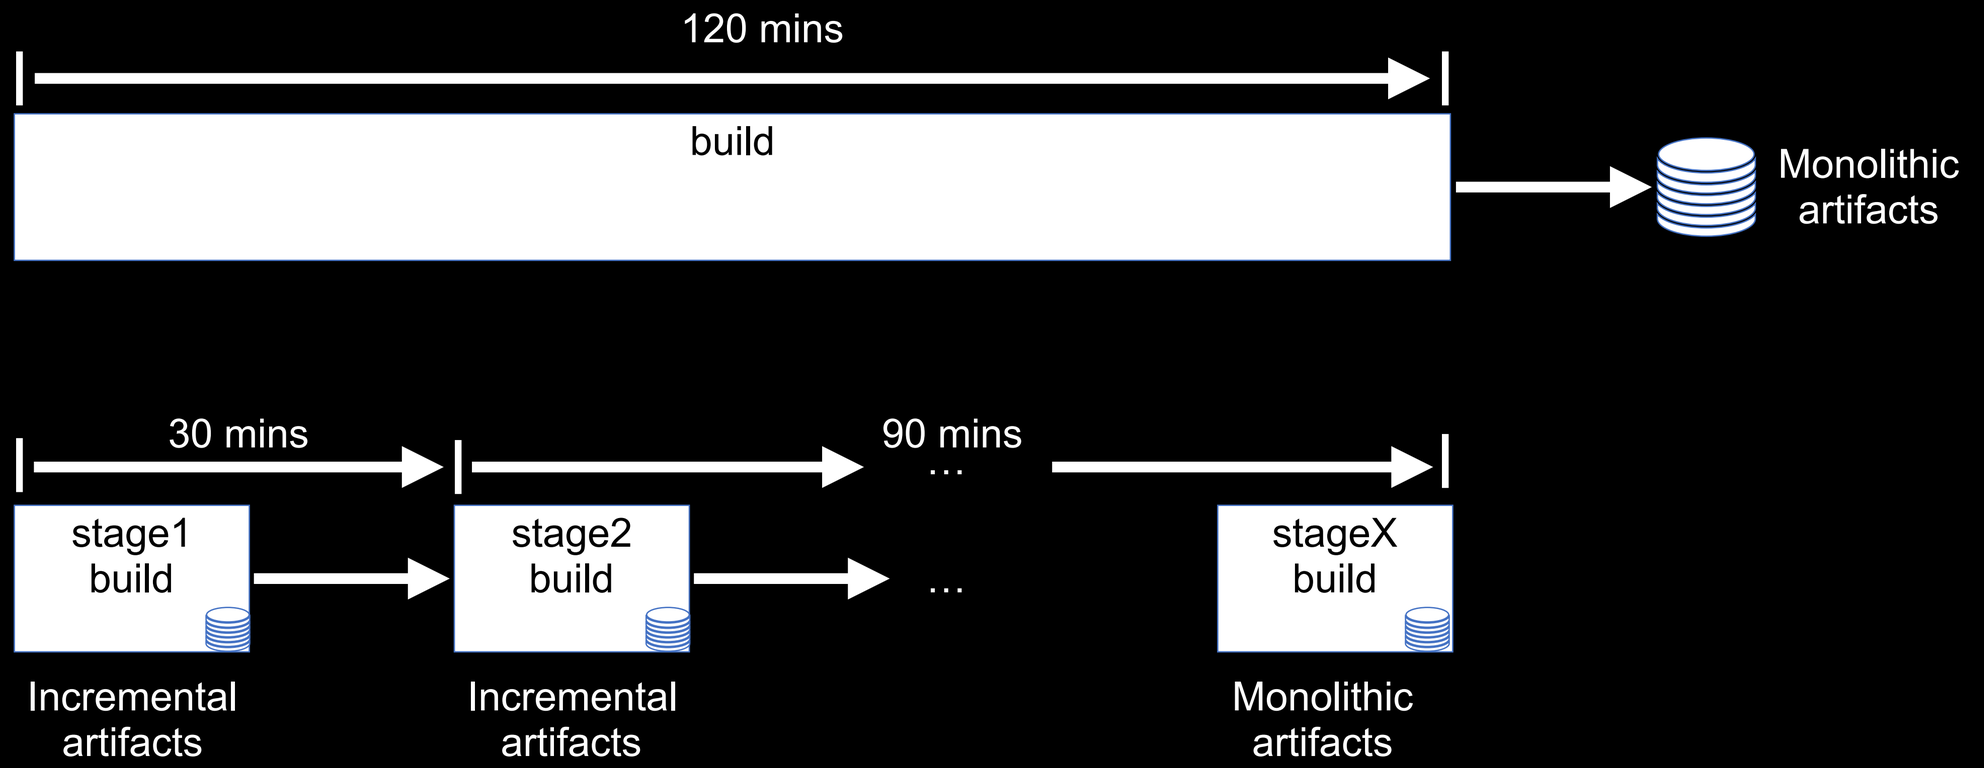

Advanced zero copy workflow

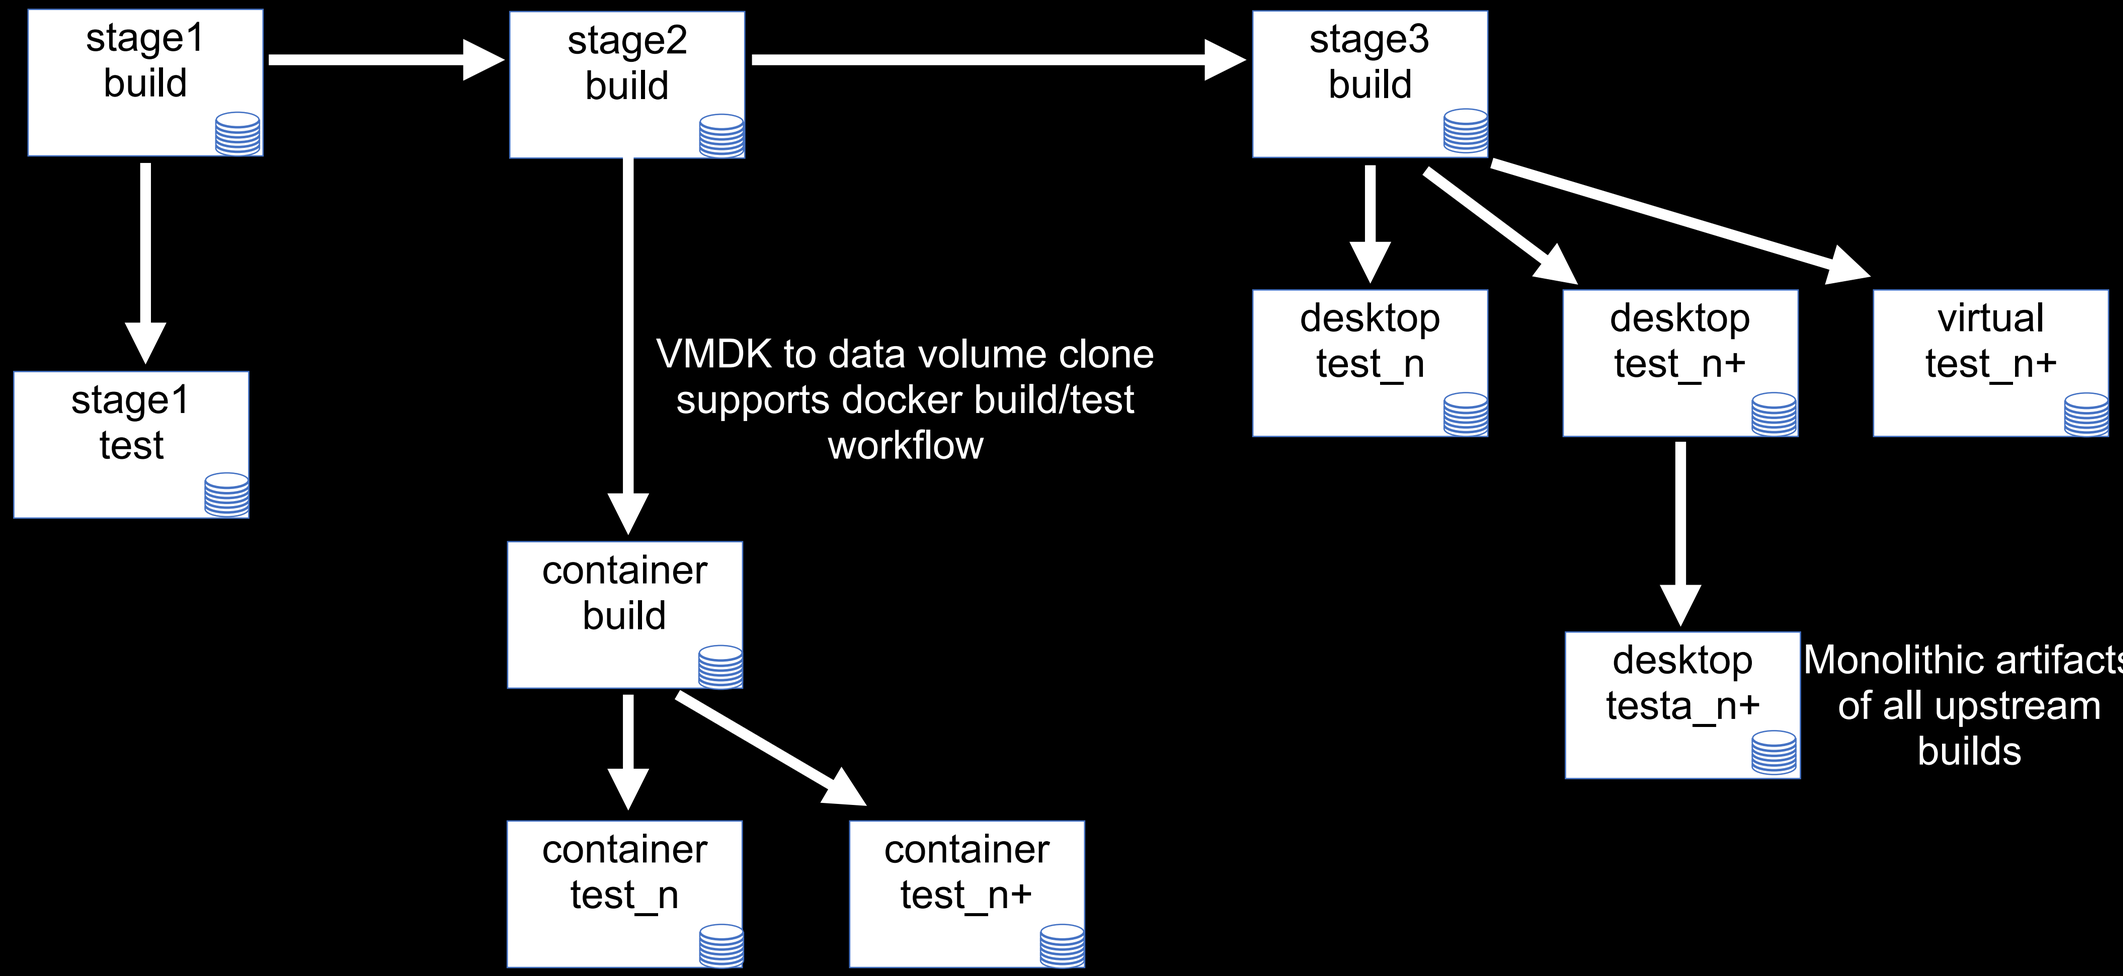

Zero copy “out-of-band” implementation

Python module utilizing Jenkins and vSphere APIs

Zero copy “out-of-band” implementation

Python module utilizing Jenkins and vSphere APIs
Runs within a build step

Zero copy “out-of-band” implementation

Python module utilizing Jenkins and vSphere APIs
Runs within a build step
Creates a “worker” VM for each build

Zero copy “out-of-band” implementation

Python module utilizing Jenkins and vSphere APIs
Runs within a build step
Creates a “worker” VM for each build
Manages artifacts

Zero copy “out-of-band” implementation

Python module utilizing Jenkins and vSphere APIs
Runs within a build step
Creates a “worker” VM for each build
Manages artifacts
In production for 3 years

Zero copy artifacts

Operates on a clone (preserves original)

Zero copy artifacts

Operates on a clone (preserves original)
Complete snapshot of entire build/test environment

Zero copy artifacts

Operates on a clone (preserves original)
Complete snapshot of entire build/test environment
Advanced developer actions (e.g., debug live session)

Zero copy artifacts

Operates on a clone (preserves original)
Complete snapshot of entire build/test environment
Advanced developer actions (e.g., debug live session)
CLI and Web browsing

Zero copy benefits

Fast and efficient

Zero copy benefits

Fast and efficient
Reduced storage consumption (data is compressed and deduplicated)

Zero copy benefits

Fast and efficient
Reduced storage consumption (data is cmpressed and deduplicated)
Advanced artifact features

Cloneable project

Based on vSphere-Cloud plugin

Cloneable project

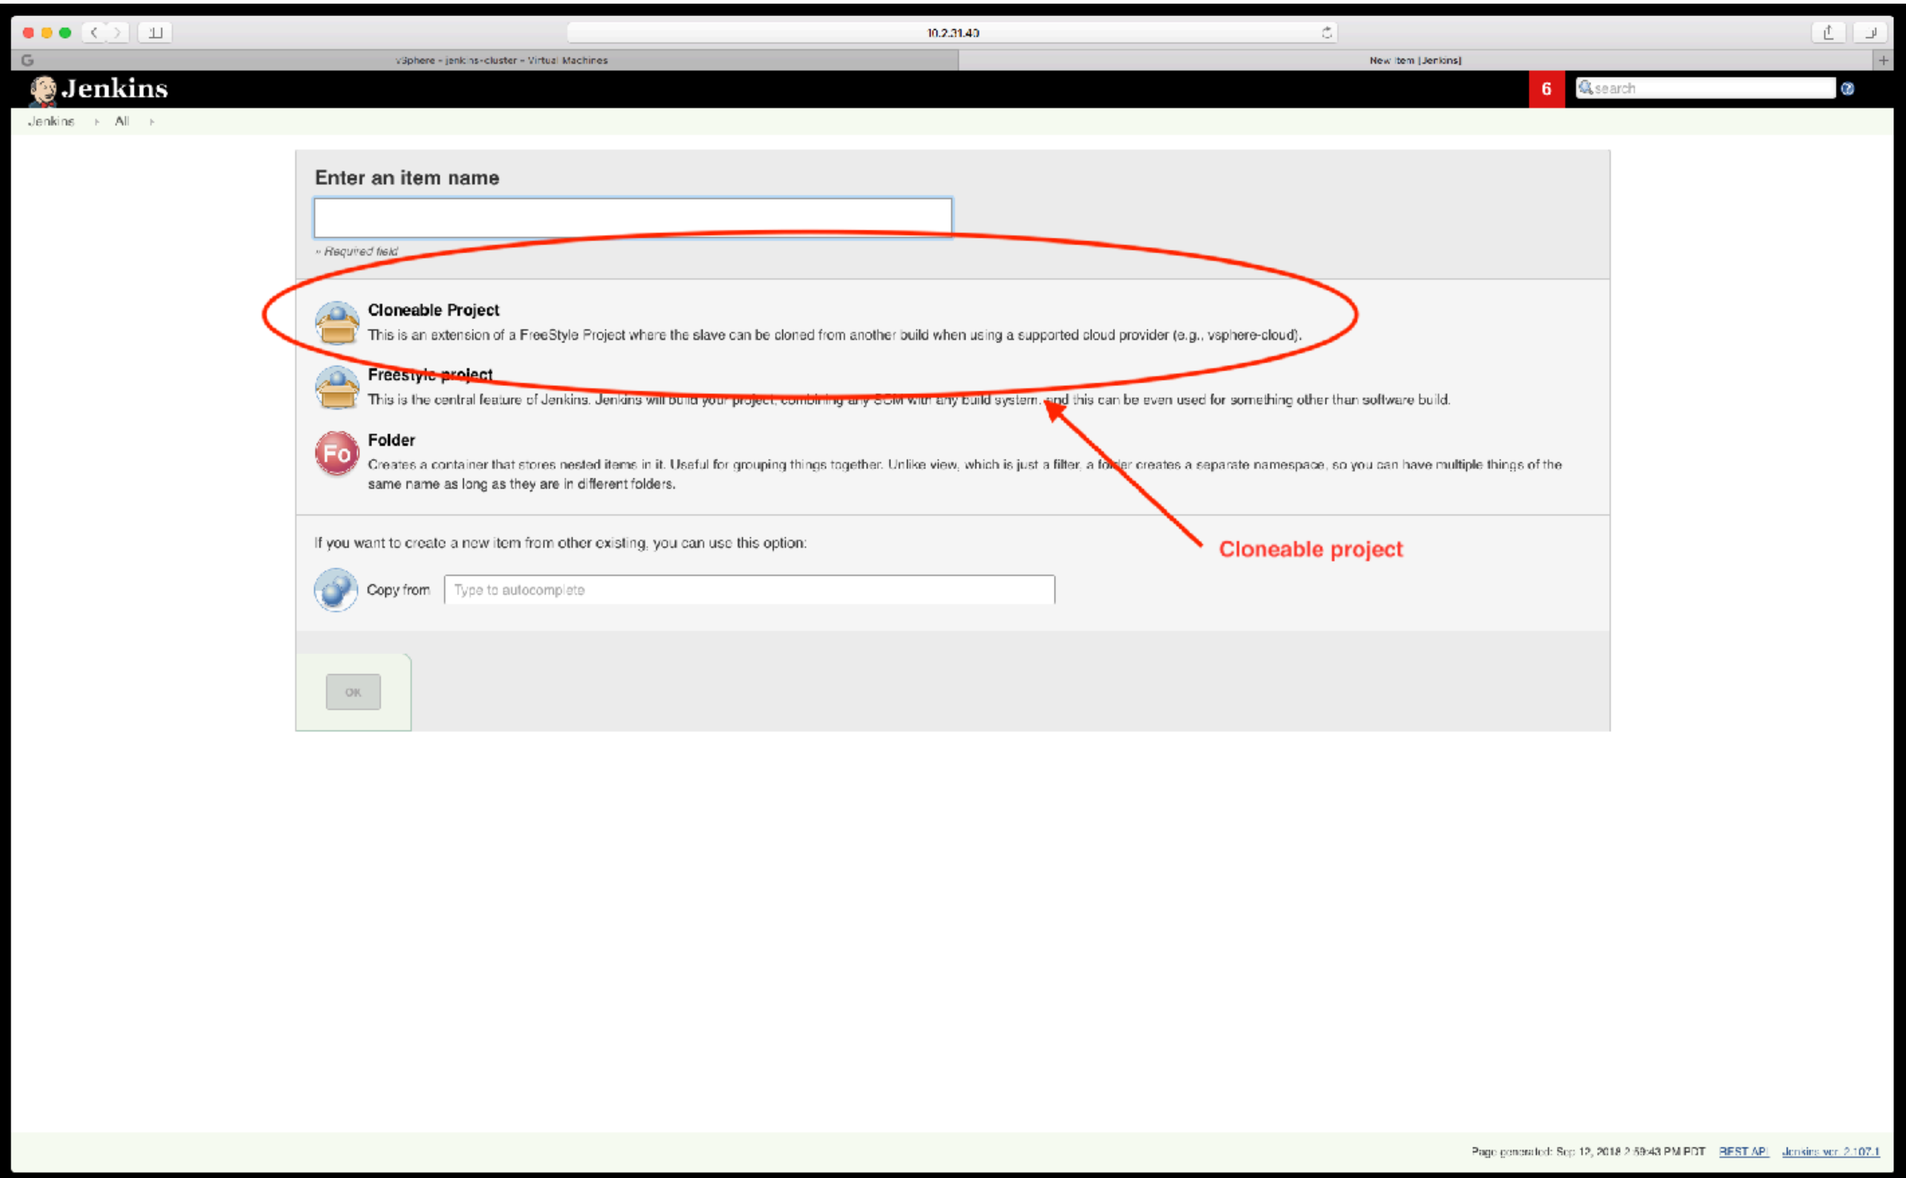

Cloneable project

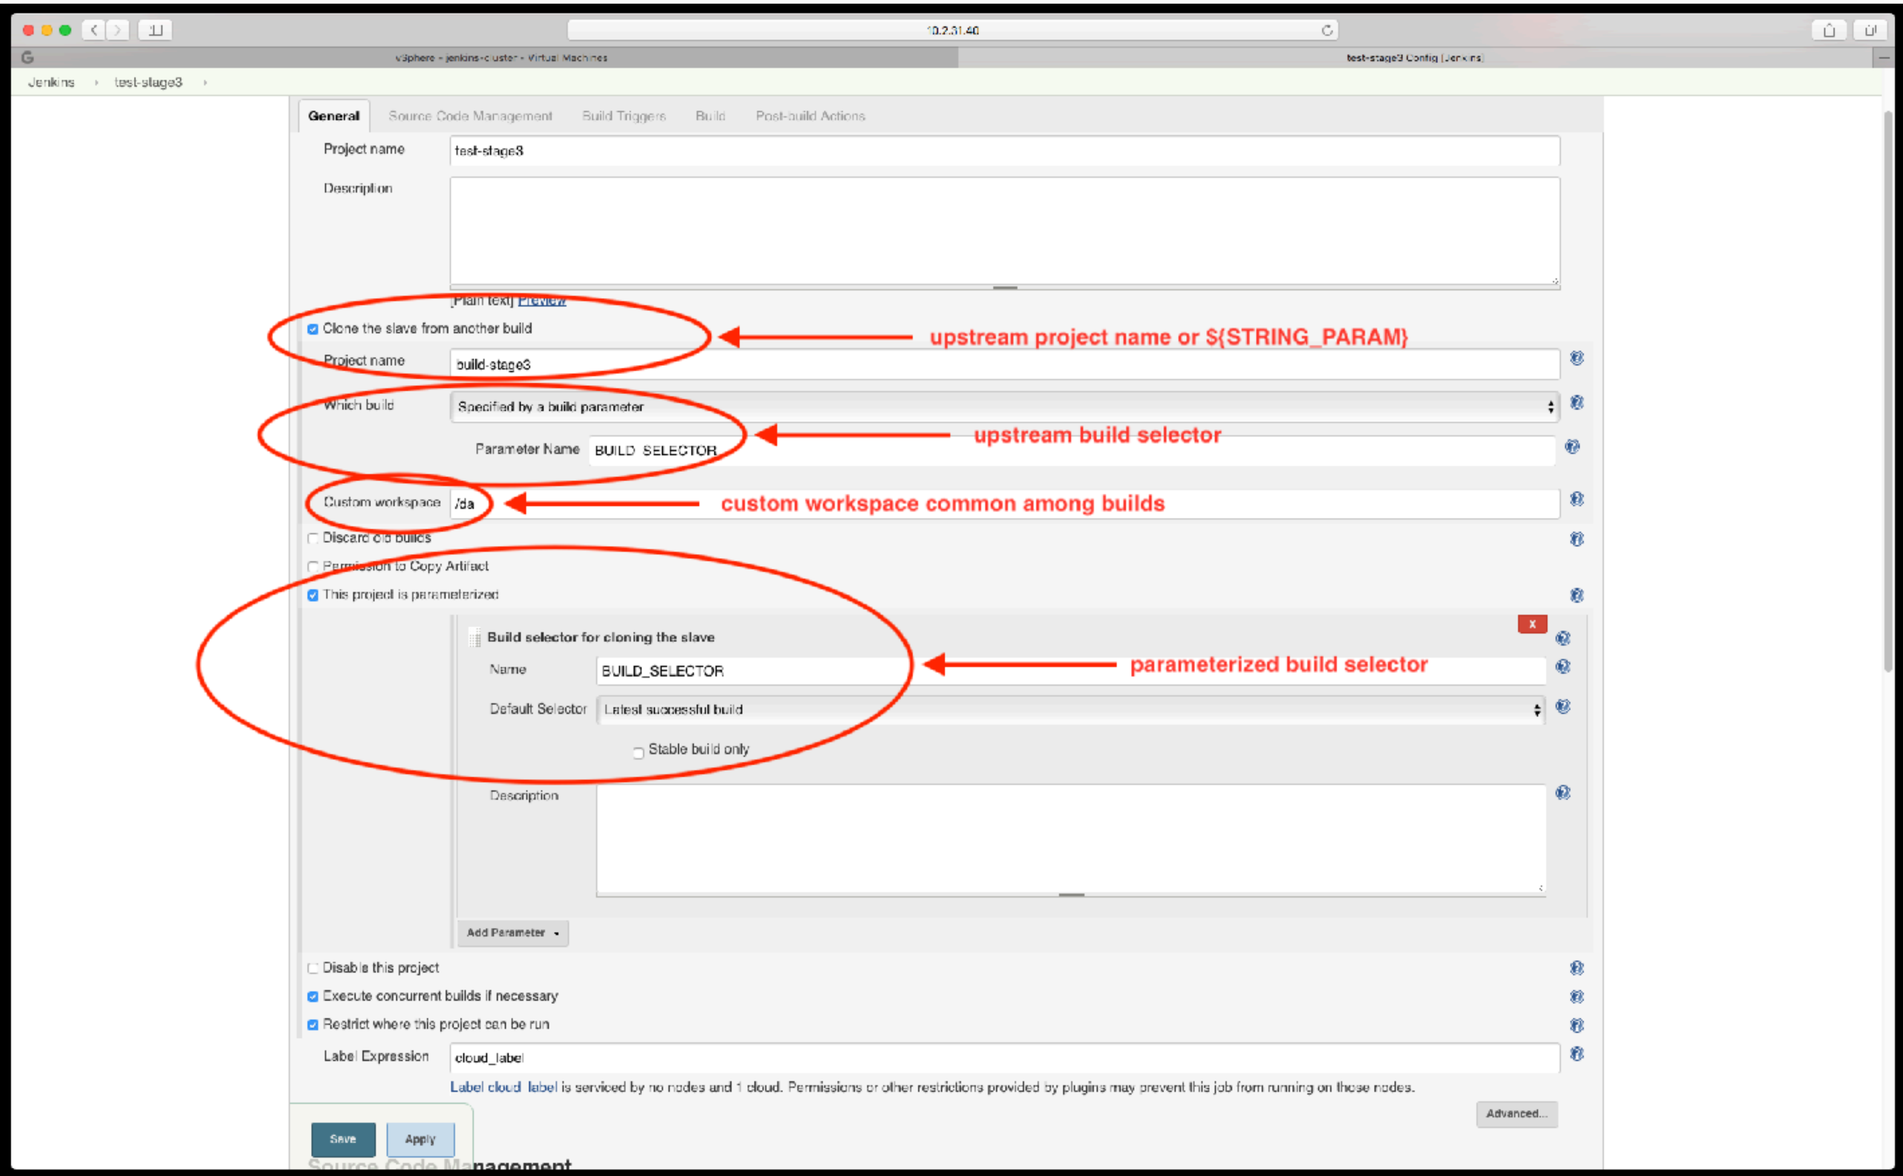

Cloneable project

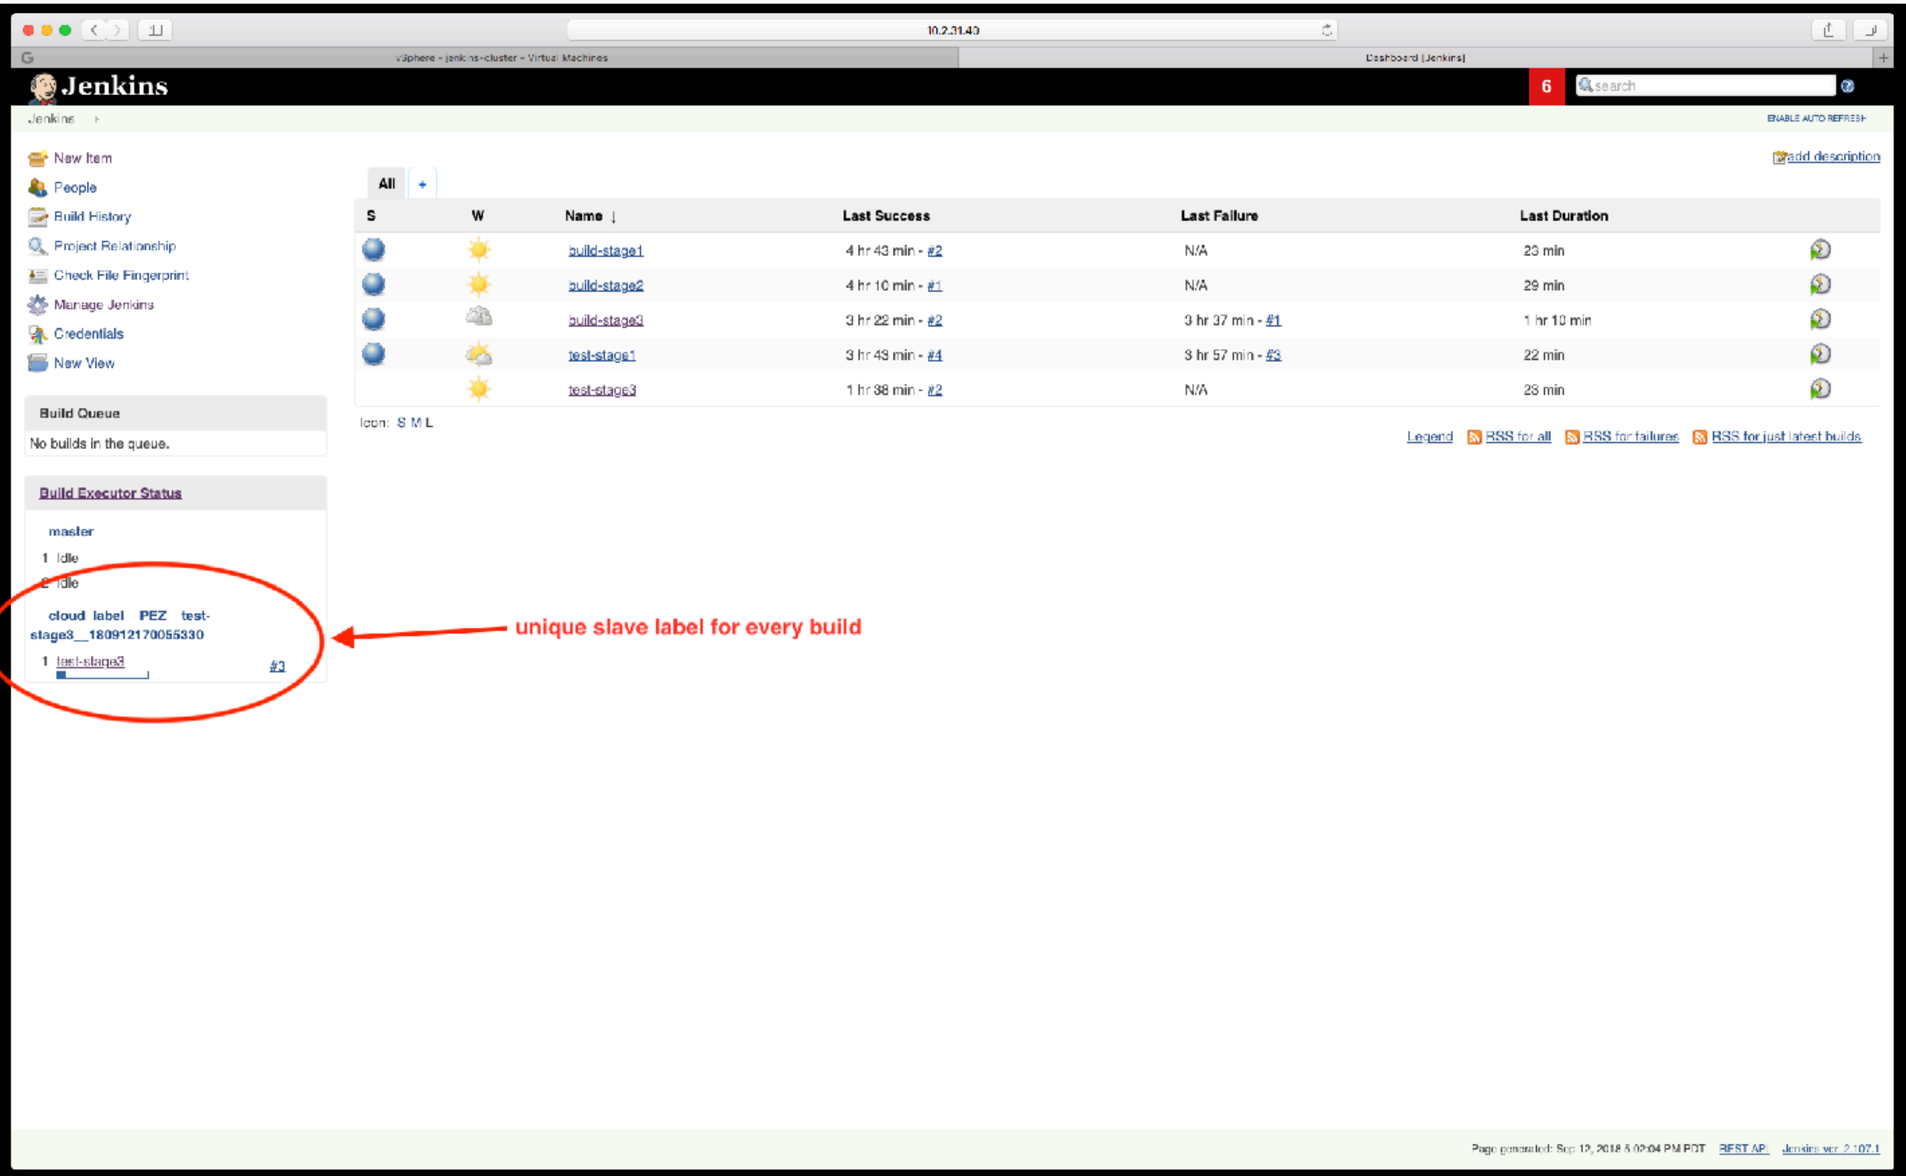

Cloneable project

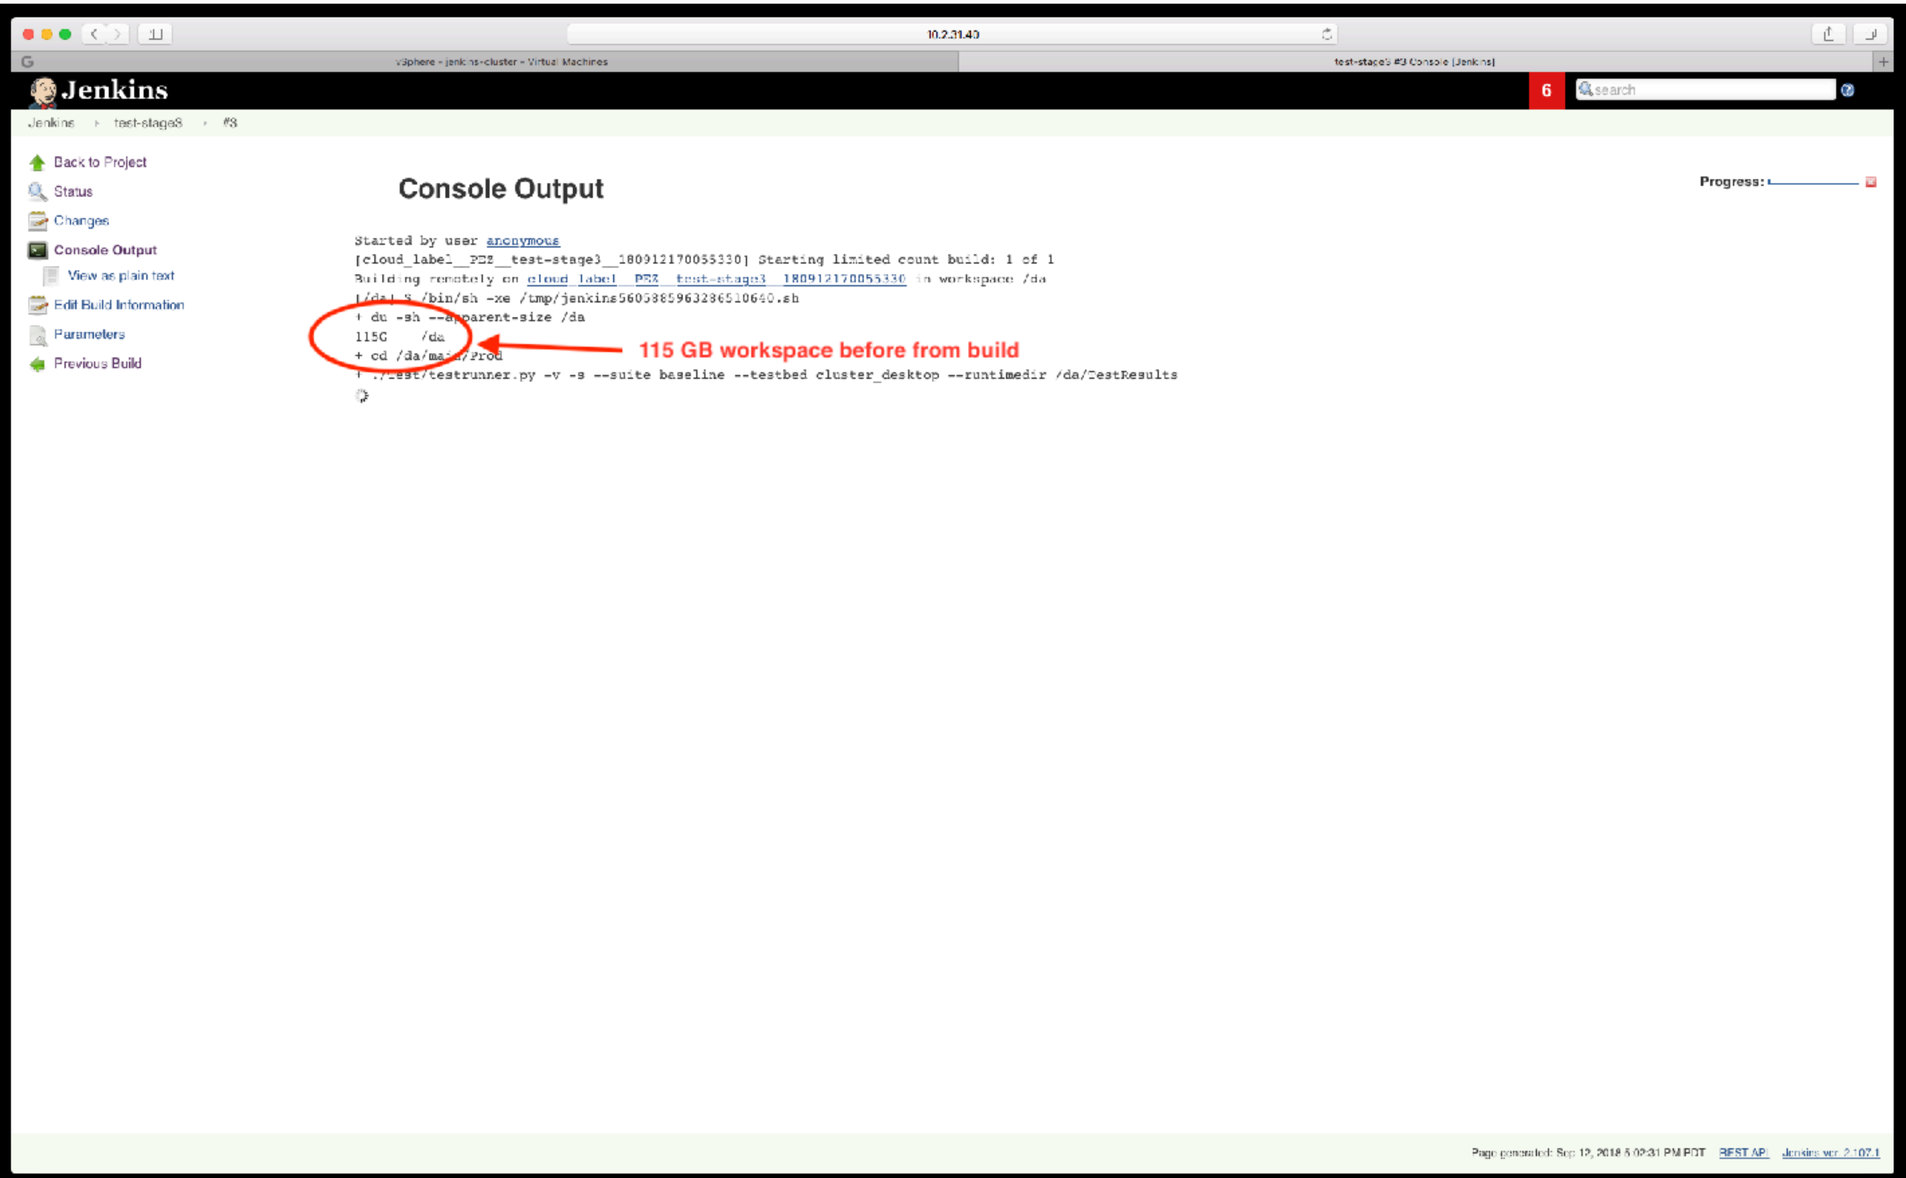

Cloneable project

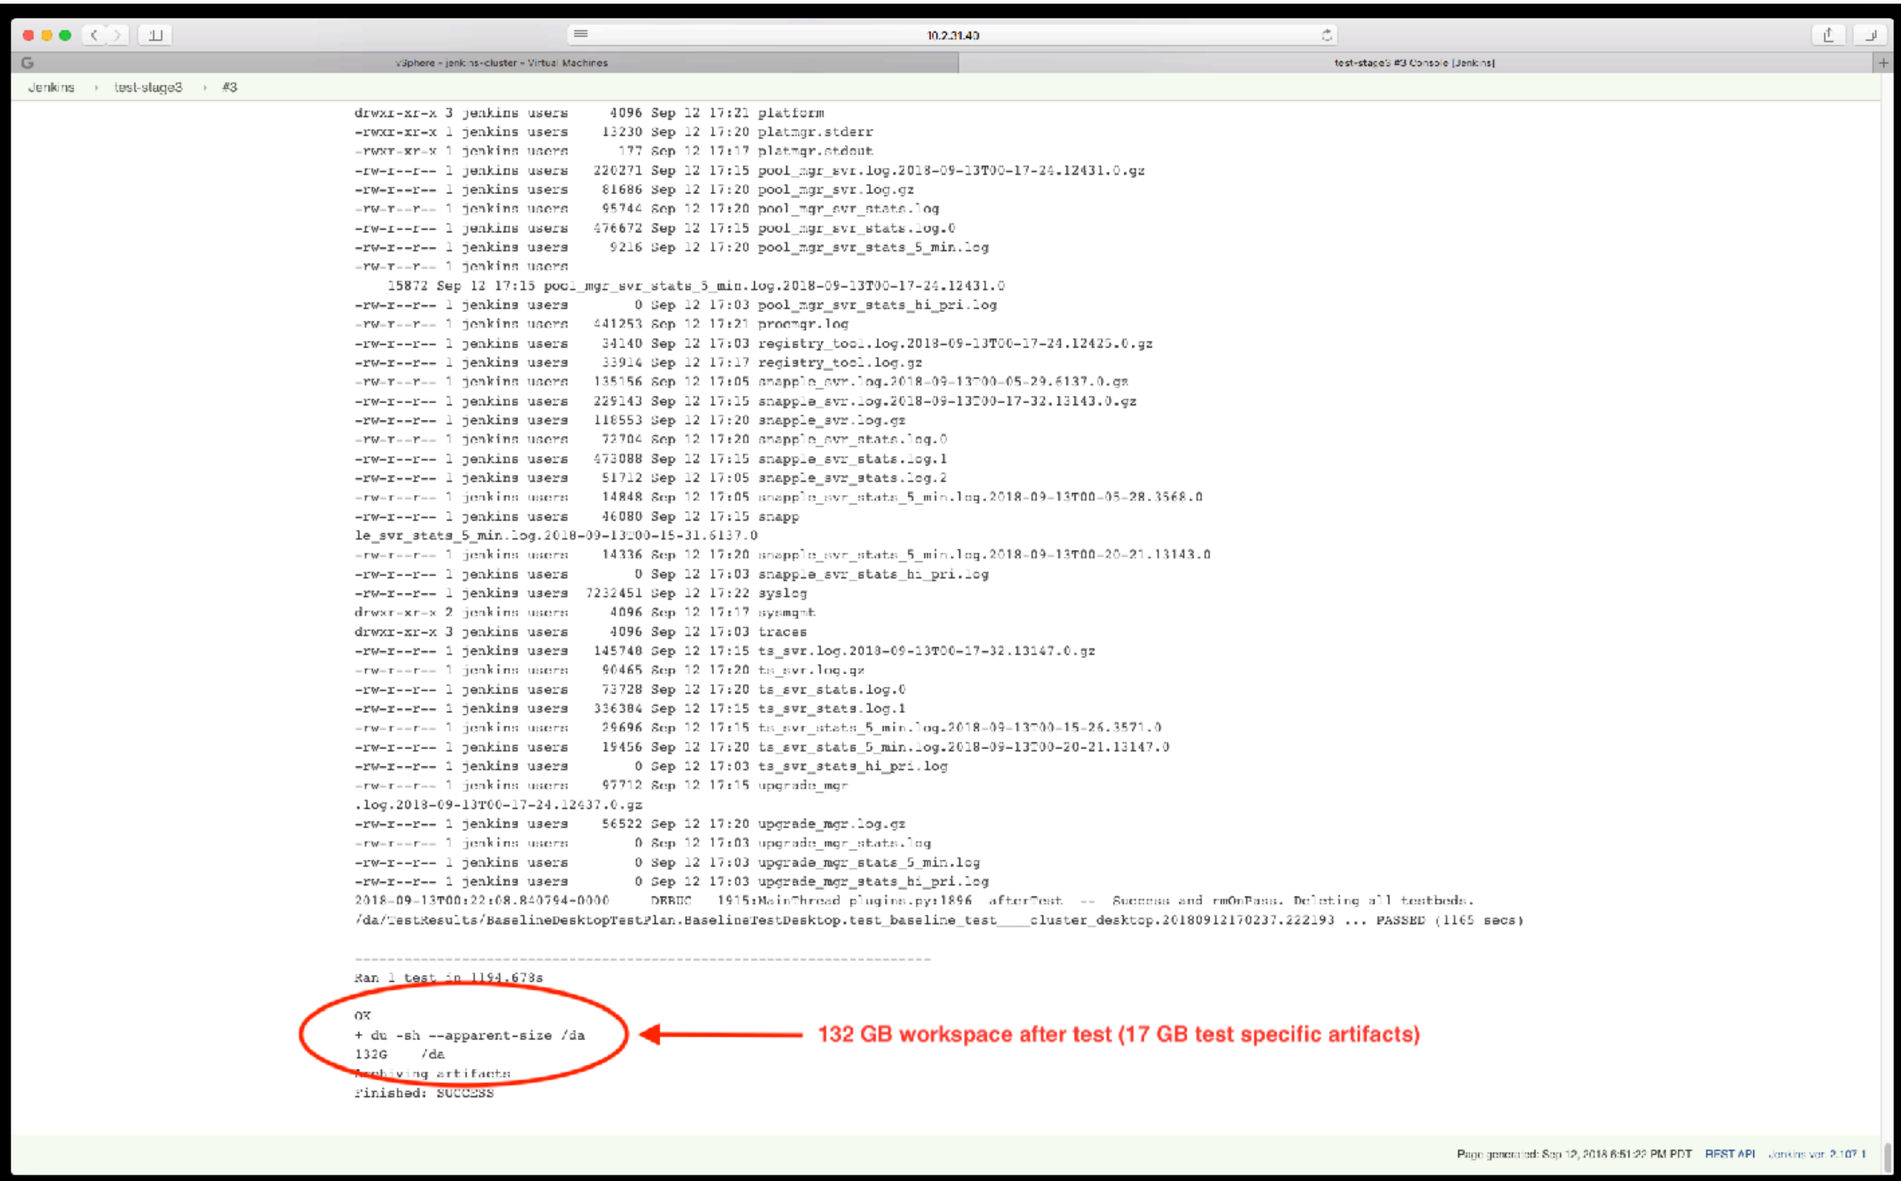

“Incremental” builds

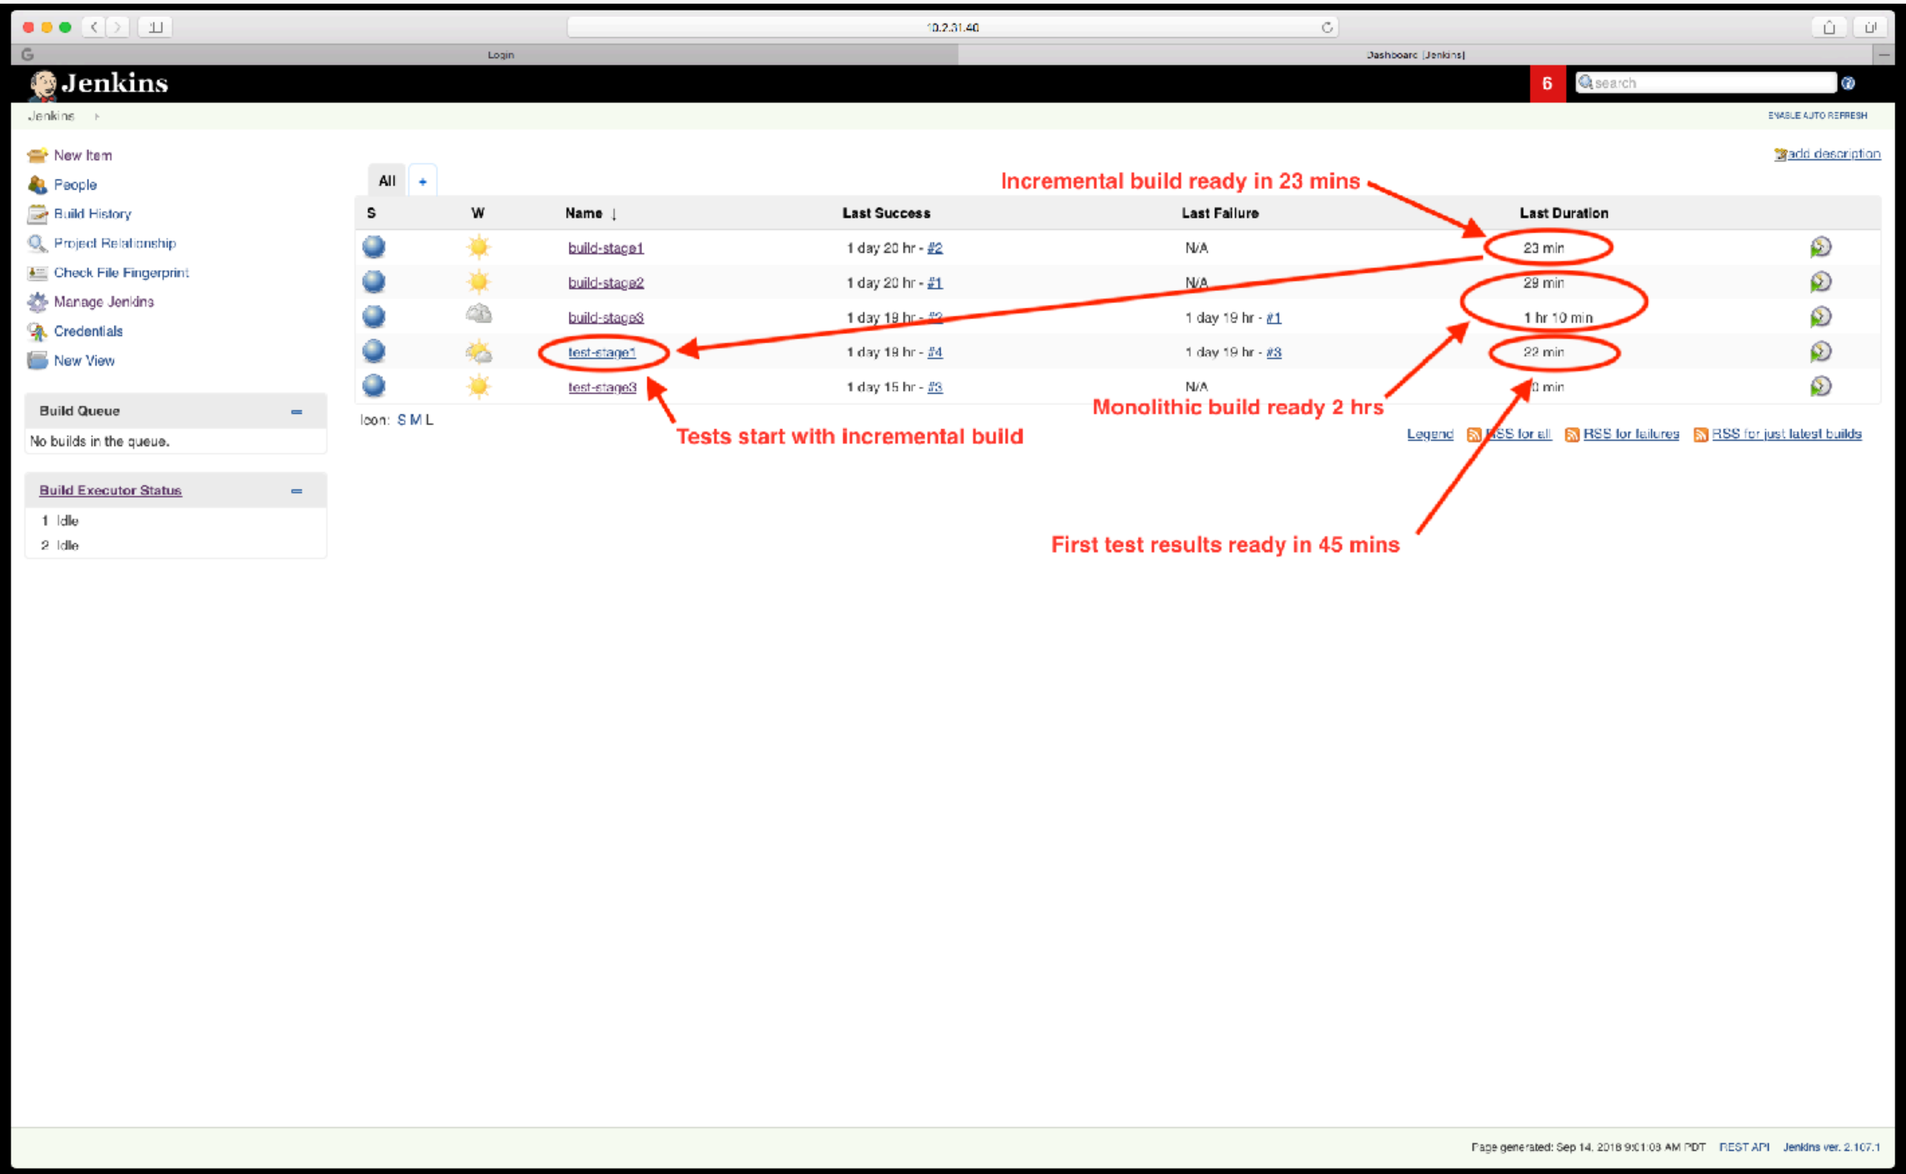

Future work

Docker-cloud plugin support (out-of-band implementation works great)

Future work

Docker-cloud plugin support (out-of-band implementation works great)
Better overall integration (artifact handling, extend cloud interface, etc.)

Thanks!

Thanks to all who help to develop and maintain Jenkins!

And for these awesome plugins that we leveraged!

Copy Artifact
vSphere-Cloud

Contacts and more info

anupam@datrium.com
kyle@datrium.com
https://github.com/datrium